

How the technology works and what its limits are

GPT Agent · PRODOC Appraisal Checklist

Presenter: Ahmed Eid

Session: Day 1 · Thursday 27 August 2026 · 9:30–10:30 (Lima)

Block: Item 4 · 30 minutes, with questions folded in

What we will cover

- 1** **Where it comes from** from the Streamlit app to the GPT Agent
- 2** **How it is built** the rubric, the tests, the search for evidence
- 3** **Why it repeats** randomness, 10 runs and stability
- 4** **What it does NOT do** real limits, and how it compares to the Enterprise GPT
- 5** **Demonstration** upload → run → read the Excel

Questions are folded into this block: interrupt me at any point.

Where it comes from: Streamlit to the GPT Agent

	Streamlit application	GPT Agent
Access	URL + environment password	A ChatGPT link
Install	Own server, dependencies	None: it opens in the browser
Learning curve	Tabs and parameters to learn	Plain conversation
Rubric	Loaded on the server	Loaded on the server (same)
Engine	The same evaluation code	The same evaluation code
Output	Downloadable Excel	Downloadable Excel (same)

The engine did not change. What changed is the front door: the Agent removes the install and interface barrier.

What a GPT is, in plain terms

- **A language model predicts text.** It has read enormous amounts of text and learned which words tend to follow others in a given context.
- **It is not a database.** It does not «look up» the PRODOC the way you search a filing system: it reads the whole thing and reasons over it there and then.
- **It is not a search engine.** There is no list of correct answers stored somewhere that the system retrieves.
- **It is a reader that follows instructions.** We give it the document, the rubric and a very precise instruction, and it returns a judgement with the quotation that supports it.

Practical consequence: the Agent can only see what is written in the PRODOC. What the specialist knows but did not document does not exist for it.

What the Agent does not do

It does not do this	Why	What it means for you
Approve or reject a PRODOC	It issues no official ILO determination	The decision stays with the appraiser
Guess what is undocumented	It only reads the text it is given	A «No» may be a documentation gap, not a design gap
Remember previous appraisals	Every run starts from zero	It does not compare across projects automatically
Rewrite the PRODOC	That is neither its purpose nor its scope	It shows you where to look, not what to write
Always give the same answer	There is inherent randomness	Which is why we repeat and measure stability

A deliberate design choice: we prefer a narrow, auditable diagnosis over a broad recommendation with no grounding in the document.

How the Appraisal Checklist was built in

1 **The rubric was digitised** each criterion became a structured row, not free text

2 **It was broken into tests** each criterion is expressed as boolean questions T1, T2, T3...

3 **A decision rule was written** an explicit formula determines Yes / Partial / No

4 **Verifiable anchors were set** terms and phrases the Agent must look for in the text

5 **Subjectivity was labelled** every criterion is tagged High / Medium / Low

The rubric lives on the server. Users never upload it: every appraisal runs against the same institutional version.

What it assesses: 76 criteria across 5 sections

Section	Content	Criteria	Subsections
1	Relevance	20	1.1 – 1.5
2	Design validity	13	2.1 – 2.4
3	Results framework and M&E	27	3.1 – 3.7
4	Implementation	14	4.1 – 4.4
5	Presentation	2	5.1 – 5.2

You can filter: the full rubric, one section («only section 3») or specific subsections («only 1.1 and 2.3»).
Filtering cuts both time and cost.

From criterion to tests: how it «reads» a criterion

A criterion is not assessed as a general impression.

It is broken into closed questions, answered one by one, each with quoted evidence:

Test	Closed question	Answer
T1	Does it distinguish the type of approach (sensitive / responsive / transformative)?	yes / no
T2	Does it articulate how the project challenges norms or power relations?	yes / no
T3	Are there dedicated actions to transform relations, not just to include women?	yes / no

DECISION: $T1 \wedge T2 \wedge T3 \rightarrow \text{Yes}$ • $T1 \vee T3$ (without all three) $\rightarrow$
Partial • $\neg T1 \wedge \neg T3 \rightarrow \text{No}$

This is what makes the result auditable: you can check, test by test, where and why the Agent reached that conclusion.

How it looks for evidence: DEDICATED vs FRAMING

A word appearing in the document is not enough. The Agent classifies each mention before counting it as evidence.

	FRAMING (does not count)	DEDICATED (does count)
What it is	The topic is mentioned, with no development of its own	The topic has its own space in the design
Examples	Lists of ≥ 3 groups • generic inclusion language • «among others» enumerations	An output or result naming it • a disaggregated indicator • a dedicated activity • a budget line • a quantified target

Hard rule: if all citable evidence is framing, the result must be «No» or «Not Found», however many times the topic is named.

A full case: from PRODOC text to verdict

WHAT THE PRODOC SAYS

§ 2.3 Gender approach

«The project adopts a gender-transformative approach, distinguishing it from merely sensitive or responsive approaches.»

§ 3.1 Beneficiaries

«Priority will be given to the participation of women, young people, persons with disabilities and indigenous peoples.»

§ 4.2 Activity 2.4

«Working sessions with cooperatives to review the credit-access criteria that exclude women titleholders. Budget: USD 18,000.»

HOW THE AGENT ASSESSES IT

✓ **T1 • Distinguishes the type of approach**
§ 2.3 names and distinguishes the three approaches.

✗ **T2 • Articulates how it challenges power norms**
it declares the approach but never explains the mechanism.

✓ **T3 • Dedicated actions to transform relations**
§ 4.2 revises credit-access rules, with its own budget.

§ 3.1 does not count as evidence
- a list of four groups with no follow-through: FRAMING, not DEDICATED.

DECISION: $T1 \wedge T2 \wedge T3$ required. T1 and T3 met, T2 missing → **PARTIAL** • illustrative extract, not a real PRODOC.

Why every criterion is run 10 times

- **A language model is not deterministic.** The same question about the same document can produce different answers.
- **Asking once would be fragile.** We would be handing you the result of a single roll.
- **The fix: repeat and consolidate.** Each criterion is evaluated 10 times independently and the most frequent result (the mode) is taken.
- **Disagreement is information, not noise.** How many of the 10 runs agreed is reported to you, and you should read it.

A full appraisal is $76 \text{ criteria} \times 10 \text{ runs} \approx \text{760 model calls}$. That is why it takes minutes, and why filtering by section pays off.

Where the randomness comes from

Source	What happens	How we control it
Probabilistic generation	The model picks among possible continuations; not always the same one	10 runs + modal result
Ambiguous criteria	Where the rubric admits a reading, the model may lean differently	Closed tests and explicit decision rules
Scattered evidence	If evidence is spread out, different runs quote different passages	The DEDICATED / FRAMING filter
Long documents	More text, more room for partial reading	Structured extraction, one criterion at a time

Randomness is not eliminated: it is measured and reported. That is better than hiding it behind a single answer.

What «stability» means

Stability (%) = of the 10 runs, how many agreed with the final result.

Institutional threshold: **80%**. Below it, the criterion is flagged for human review.

Stability	Reading	What to do
100 – 80%	The model was consistent	Treat as a solid diagnosis; still verify the evidence
79 – 50%	The runs disagreed	Human review required; read the «Resultado Alternativo»
< 50%	No real consensus	The criterion is ambiguous or the evidence thin: you decide

Stability does not measure whether the Agent was right. It measures how confident it was internally. A confidently wrong 100% is possible: always verify the evidence.

Subjective criteria and human review

- Every criterion carries a subjectivity label: **High • Medium • Low**.
- **High subjectivity** means the judgement depends on institutional context, not just the text. Example: whether a gender approach is «transformative».
- **These criteria get more reasoning** from the model and are always flagged in the «Revisión humana recomendada» column.
- **Two triggers for that flag:** high subjectivity, or stability below 80%.

Use it as a work queue: start with the flagged criteria. That is where your professional judgement adds the most.

Strengths and trade-offs of this technology

	Strength	Trade-off / cost
Coverage	Reviews all 76 criteria without fatigue	It reviews what is written, not what is known
Time	Frees specialist hours to work on the gaps	A full run takes several minutes
Traceability	Quotes the evidence and shows the reasoning per test	It produces a lot of detail: you must know how to read it
Consistency	Applies the same rubric to every project	Not deterministic: hence the repeats
Cost	Low per document against the equivalent human time	Billed per call: the budget is finite

Tokens and budget: cost depends on the size of the PRODOC and how many criteria are run. Filtering by section is the direct saving lever.

What about the ILO Enterprise GPT?

	Enterprise GPT	Appraisal Checklist Agent
Rubric	Would have to be uploaded in every conversation	Server-side, single version
Application	General reading of the attached file	Each criterion run separately, with its tests
Repetition	One single pass	10 runs per criterion + modal result
Stability	Not available	Reported per criterion, 80% threshold
Output	Text in the conversation	Structured Excel, auditable and archivable
Access today	No access to the Agent from the enterprise environment	A ChatGPT link

The added value is not «using AI»: it is the institutional rubric applied criterion by criterion, repeated and measured. A generic GPT does not reproduce that.

The three limits to keep in mind

1 It is a **complete first review**, not a deterministic diagnosis. It flags gaps for you to examine.

2 It works on the **document's evidence**, not on the real project or on what you know about it.

3 It does not **replace the appraiser's judgement**. The result is an input; the determination stays human.

No output constitutes an official ILO determination. It is an AI-assisted assessment that requires expert validation.

Part 2 • The tool in use

Upload → scope → run → download → interpret

The full flow, in five steps

- 1** **Upload the PRODOC** one .docx file per appraisal
- 2** **Set the scope** full rubric, one section, or specific subsections
- 3** **Run** the Agent starts the job and tracks it until it finishes
- 4** **Download the Excel** it is the auditable record of the appraisal
- 5** **Interpret** locate the evidence and prioritise the flagged criteria

State the scope in the same message as the upload: it saves a round of questions and caps the cost from the start.

Steps 1 and 2 • Upload and scope

- **Open the Agent and say hello.** It introduces itself: what it assesses, which sections exist, what you can filter. That costs no appraisal.
- **Attach a single .docx.** If you upload several, the Agent will ask which one: one document per appraisal.
- **Say what you want assessed.** Example instructions:
 - «Assess this PRODOC against the full rubric»
 - «Assess only section 3 (results framework) and summarise the gaps»
 - «Assess only 1.1 and 2.3»

Where to start: one section first. It arrives sooner, costs less, and lets you calibrate your reading before a full run.

Step 3 • What happens while it runs

1 The Agent downloads the document and extracts its text by section

2 It loads the server-side rubric and filters to the criteria you asked for

3 It runs the evaluations in parallel each criterion, 10 times, up to 48 concurrent calls

4 It consolidates each criterion modal result + stability percentage

5 It builds the Excel and attaches it to the conversation

No need to ask «is it done yet?»: the Agent reports progress on its own — «180/760 (24%), about 3 minutes left» — until it delivers the file.

Step 4 • The Excel: sheet «Resultado Diagnostico»

A single sheet, one row per criterion assessed. It is the auditable record: keep it.

Column group	What it holds	What it is for
Identification	ID • Subsección • Criterio • Transversales	Locate the criterion in the Checklist
Result	Respuesta (Yes / Partial / No / Not Found / N/A)	The diagnosis for that criterion
Confidence	Estabilidad (%) • Estable ($\geq 80\%$) • Resultado Alternativo	How far the 10 runs agreed
Support	Razonamiento (check by check) • Evidencia	Verify why it reached that result
Priority	Revisión humana recomendada	Your work queue

Column headers are in Spanish, as the file produces them. «Not Found» is not «No»: «No» says the criterion is not met; «Not Found» says the document does not let you tell.

How to read a row, in order

- 1 Look at the Respuesta** Yes / Partial / No / Not Found / N/A
- 2 Look at the Estabilidad** did the 10 runs agree, or was there disagreement?
- 3 Read the Evidencia** does that passage genuinely support the result?
- 4 Read the Razonamiento** which check failed, and on what grounds
- 5 Decide** design gap, documentation gap, or Agent error?

Never accept an answer without opening the evidence. The evidence is what turns an automated result into a defensible diagnosis.

What the «Razonamiento» column contains

POR QUÉ PARTIAL · Se cumplen 2 de 3 chequeos.

Falta: articular cómo el proyecto cuestiona normas de poder.

VERIFICACIÓN

✓ ¿Distingue el tipo de enfoque?

la sección 3.2 distingue los tres enfoques.

✗ ¿Articula cómo cuestiona normas de poder?

no se explica el mecanismo en ninguna sección.

✓ ¿Acciones dedicadas a transformar relaciones?

actividad 2.4, con presupuesto asignado.

ESTABILIDAD · 6 de 10 corridas coincidieron. Alternativo: No.

REGLA · se requieren los 3 chequeos ($T1 \wedge T2 \wedge T3$)

POR QUÉ / why

the reason, in one line. If you are in a hurry, read only this.

VERIFICACIÓN / checks

each check stated in full, with ✓ / ✗ and its justification.

ESTABILIDAD / stability

how many of the 10 runs agreed.

REGLA / rule

the formal Checklist rule, for audit.

No more decoding « $T1 \wedge T2 \wedge T3$ »: each check reads on its own. The output is in Spanish; the formal rule stays at the foot for anyone tracing the result.

Step 5 • Locating the evidence in the PRODOC

- **The Evidencia column holds the verbatim quotation** the Agent used. Search for it in the PRODOC with Ctrl+F to see it in context.
- **When the evidence is an absence**, the Agent says so explicitly: «No se encontró sección X». That is a verifiable finding too.
- **Three possible outcomes when you verify:**
 - The information is there and the Agent missed it** → discard the automated diagnosis
 - It is not written, but the formulator knows it** → improve the PRODOC
 - It does not exist yet** → resolve before closing formulation, or carry into inception

The formula: does it exist? → is it sufficient? → is it documented?

Linking to other ILO resources

Resource	What it adds	Example use
Chat with the Agent	Ask about the result already produced, in plain language	«Summarise the section 3 gaps in three points»
The other GPTs built	Specific Attributes (gender, participation, just transition) and Sustainability	Go deeper on a theme the Checklist only touches transversally
i-EVal chatbot	Evidence from previous evaluations in the region	«What recurring problems have evaluations found in projects relying on labour ministries?»

The Excel is not the end of the process: it is the starting point of the conversation with the formulating team.

Frequently asked questions

Question	Short answer
Why can an answer change?	The model is not deterministic. Hence 10 repeats and a reported stability figure.
What does 50–70% stability mean?	The runs disagreed: ambiguous criterion or thin evidence. Human review.
Why did it miss something I know is there?	It only reads the document. If it is there and was missed, verify and discard the diagnosis.
Why no improvement recommendations?	It holds no knowledge beyond the document. Recommendations come from your reading.
What does a PRODOC cost to assess?	It depends on document size and how many criteria are run. Filtering cuts the cost.
Can I assess other documents?	Dedicated agents exist for gender, participation, just transition and sustainability.

What to take away from this block

- **The Agent reads the document, not the project.** A «No» may be a design gap or simply something undocumented. Telling them apart is your job.
- **It runs 10 times and tells you how far they agreed.** Stability is a measure of internal confidence, not of correctness.
- **Every result comes with its evidence.** Always verify it: that is what makes the diagnosis defensible to the formulating team.
- **The «Revisión humana recomendada» column is your work queue.** Start there.

The value of the tool depends on the quality of the human review that follows.

Questions

GPT Agent · PRODOC Appraisal Checklist · Ahmed Eid